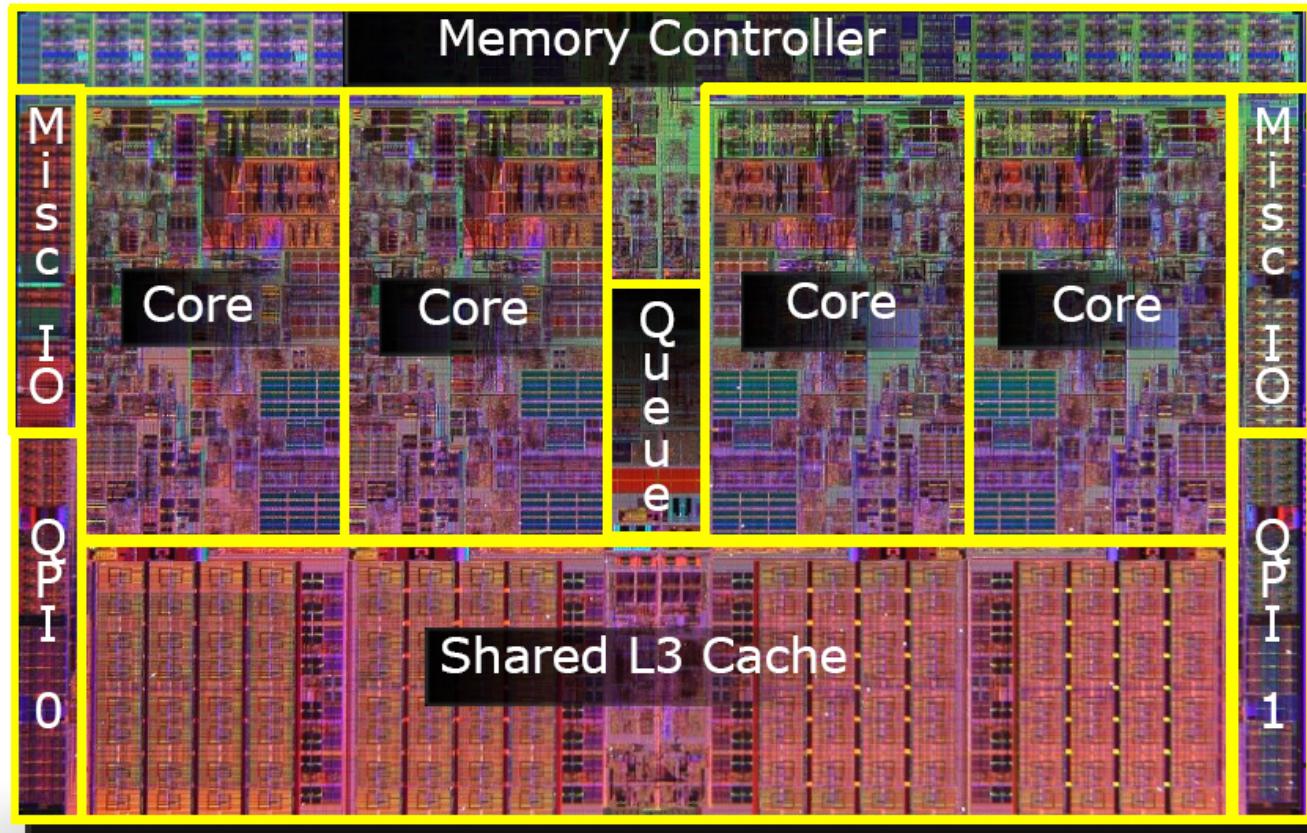


06_03_locality_and_memory hierarchy

The First Nehalem Processor



QPI: Intel® QuickPath
Interconnect

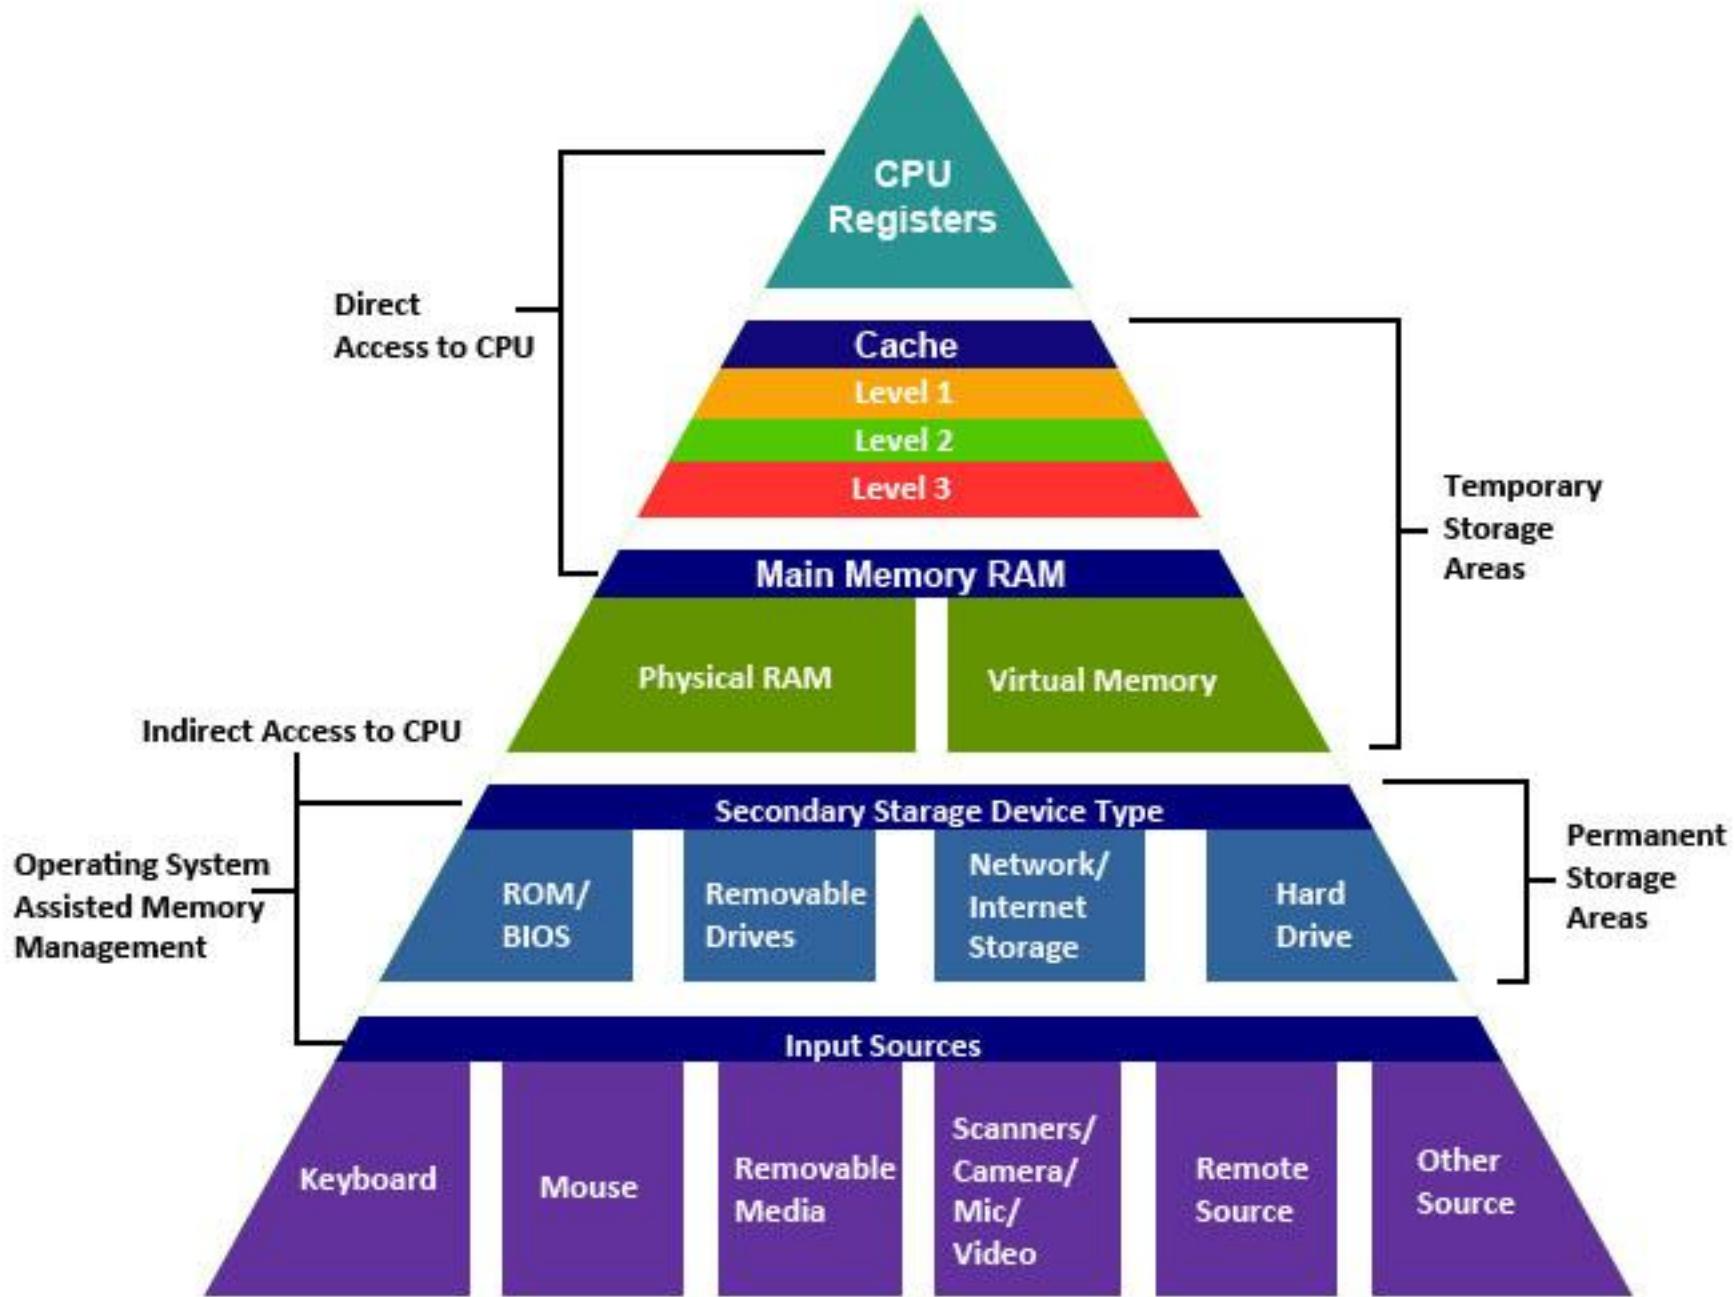
Intel Developer
FORUM

A Modular Design for Flexibility

Nehalem: Next Generation Intel® Microarchitecture 1



From [Intel \(\[http://download.intel.com/pressroom/kits/corei7/images/Nehalem_Die_callout.jpg\]\(http://download.intel.com/pressroom/kits/corei7/images/Nehalem_Die_callout.jpg\)\)](http://download.intel.com/pressroom/kits/corei7/images/Nehalem_Die_callout.jpg)



Example Memory Hierarchy

Smaller,
faster,
and
costlier
(per byte)
storage
devices

Larger,
slower,
and
cheaper
(per byte)
storage
devices

L6:

Remote secondary storage
(e.g., Web servers)

L5:

Local secondary storage
(local disks)

L2:

L2 cache
(SRAM)

L3 cache
(SRAM)

Main memory
(DRAM)

L0:

Regs

CPU registers hold words
retrieved from the L1 cache.

L1 cache holds cache lines
retrieved from the L2 cache.

L2 cache holds cache lines
retrieved from L3 cache

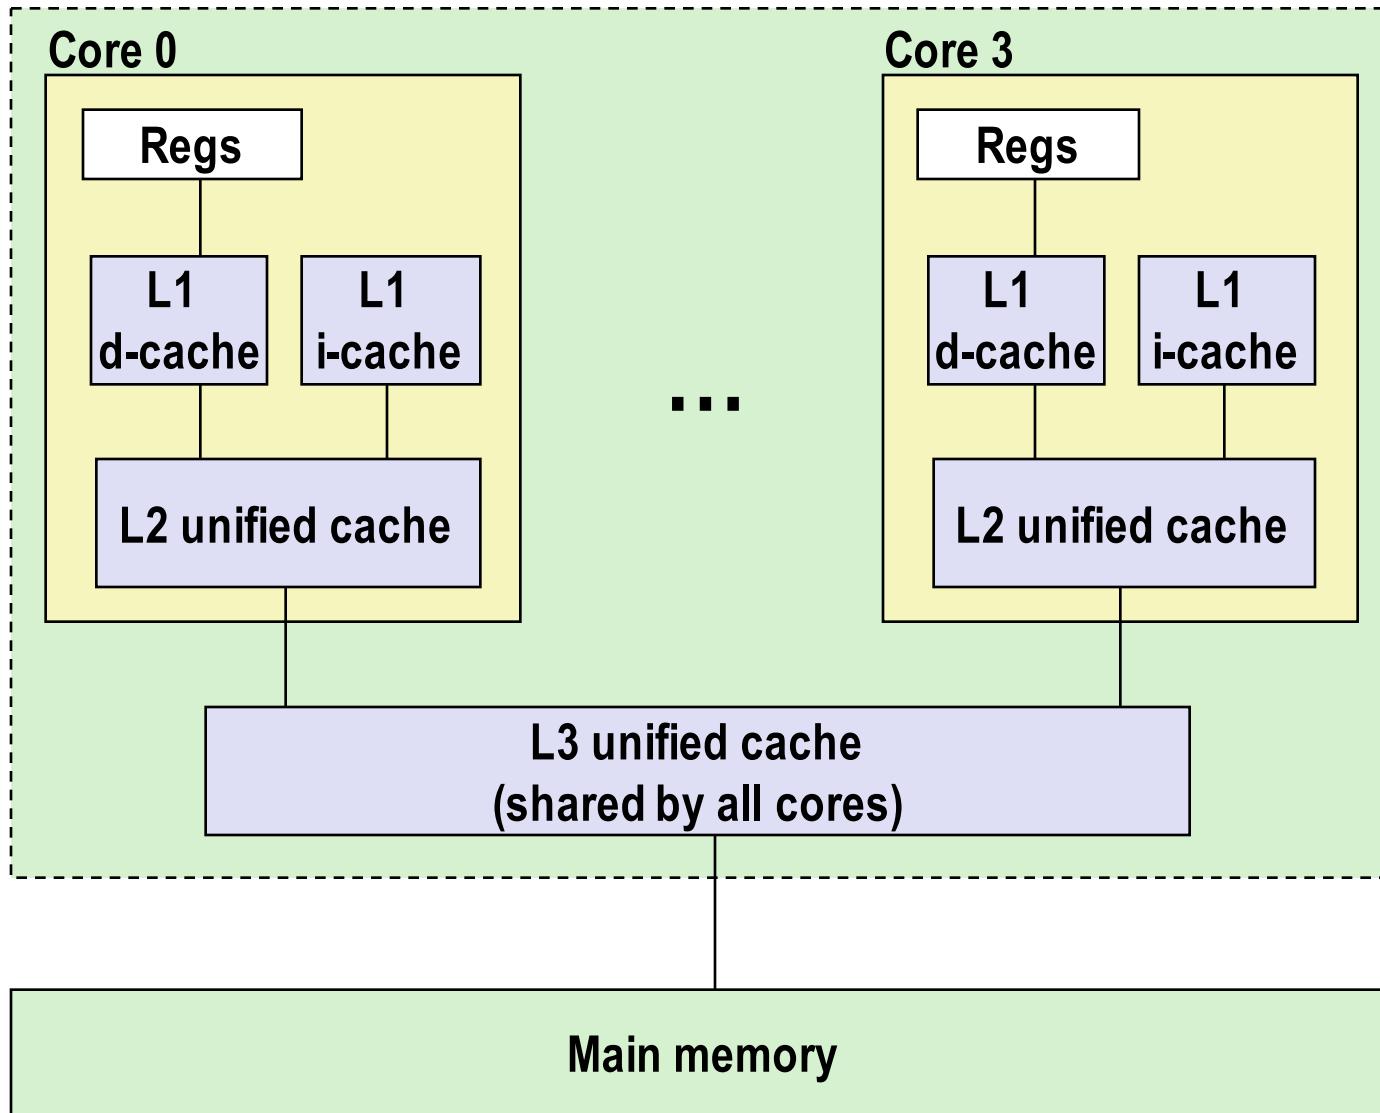
L3 cache holds cache lines
retrieved from main memory.

Main memory holds
disk blocks retrieved
from local disks.

Local disks hold files
retrieved from disks
on remote servers

Intel Core i7 Cache Hierarchy

Processor package



L1 i-cache and d-cache:
32 KB, 8-way,
Access: 4 cycles

L2 unified cache:
256 KB, 8-way,
Access: 10 cycles

L3 unified cache:
8 MB, 16-way,
Access: 40-75 cycles

Block size: 64 bytes for all caches.

Numbers Everyone Should Know

Jeff Dean talk at Stanford

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	100 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	10,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from network	10,000,000 ns
Read 1 MB sequentially from disk	30,000,000 ns
Send packet CA → Netherlands→CA	150,000,000 ns

Fake Problem

location	access time
L1 cache	0.5 ns
L2 cache	7 ns
RAM	100 ns
hard drive	10 ms
DVD	140 ms

Pretend for a minute

location	access time
L1 cache	0.5 ns
L2 cache	7 ns
RAM	100 ns
hard drive	10 ms
DVD	140 ms

location	<i>fake access time</i>
L1 cache	1 sec
L2 cache	
RAM	
hard drive	
DVD	

If we keep the ratios the same as on the LHS,
what are the remaining numbers on the RHS?

Pretend for a minute

location	access time
L1 cache	0.5 ns
L2 cache	7 ns
RAM	100 ns
hard drive	10 ms
DVD	140 ms

location	<i>fake access time</i>
L1 cache	1 sec
L2 cache	14 sec
RAM	200 sec or 3 min, 20 sec
hard drive	20,000,000 sec or \approx 231.5 days
DVD	280,000,000 sec or \approx 8.9 years

Pretend again. Let's bake a cake.



if the counter is the CPU and the L1 cache is the cabinet two feet above ...



location	<i>fake distance</i>
L1 cache	2 feet
L2 cache	28 feet
RAM	400 feet
hard drive	40,000,000 feet or about 7,500 miles about twice the distance from Philly to Nome, AK
DVD	560,000,000 feet or about 106,000 miles or about the distance around the earth 4.25 times

The moral of the story

- If you're baking a cake and you have to walk to Nome, AK to get the flour:
 1. get what you need for the rest of the recipe
 2. pick up some sugar

Locality

- Temporal locality
- Spatial locality

locality

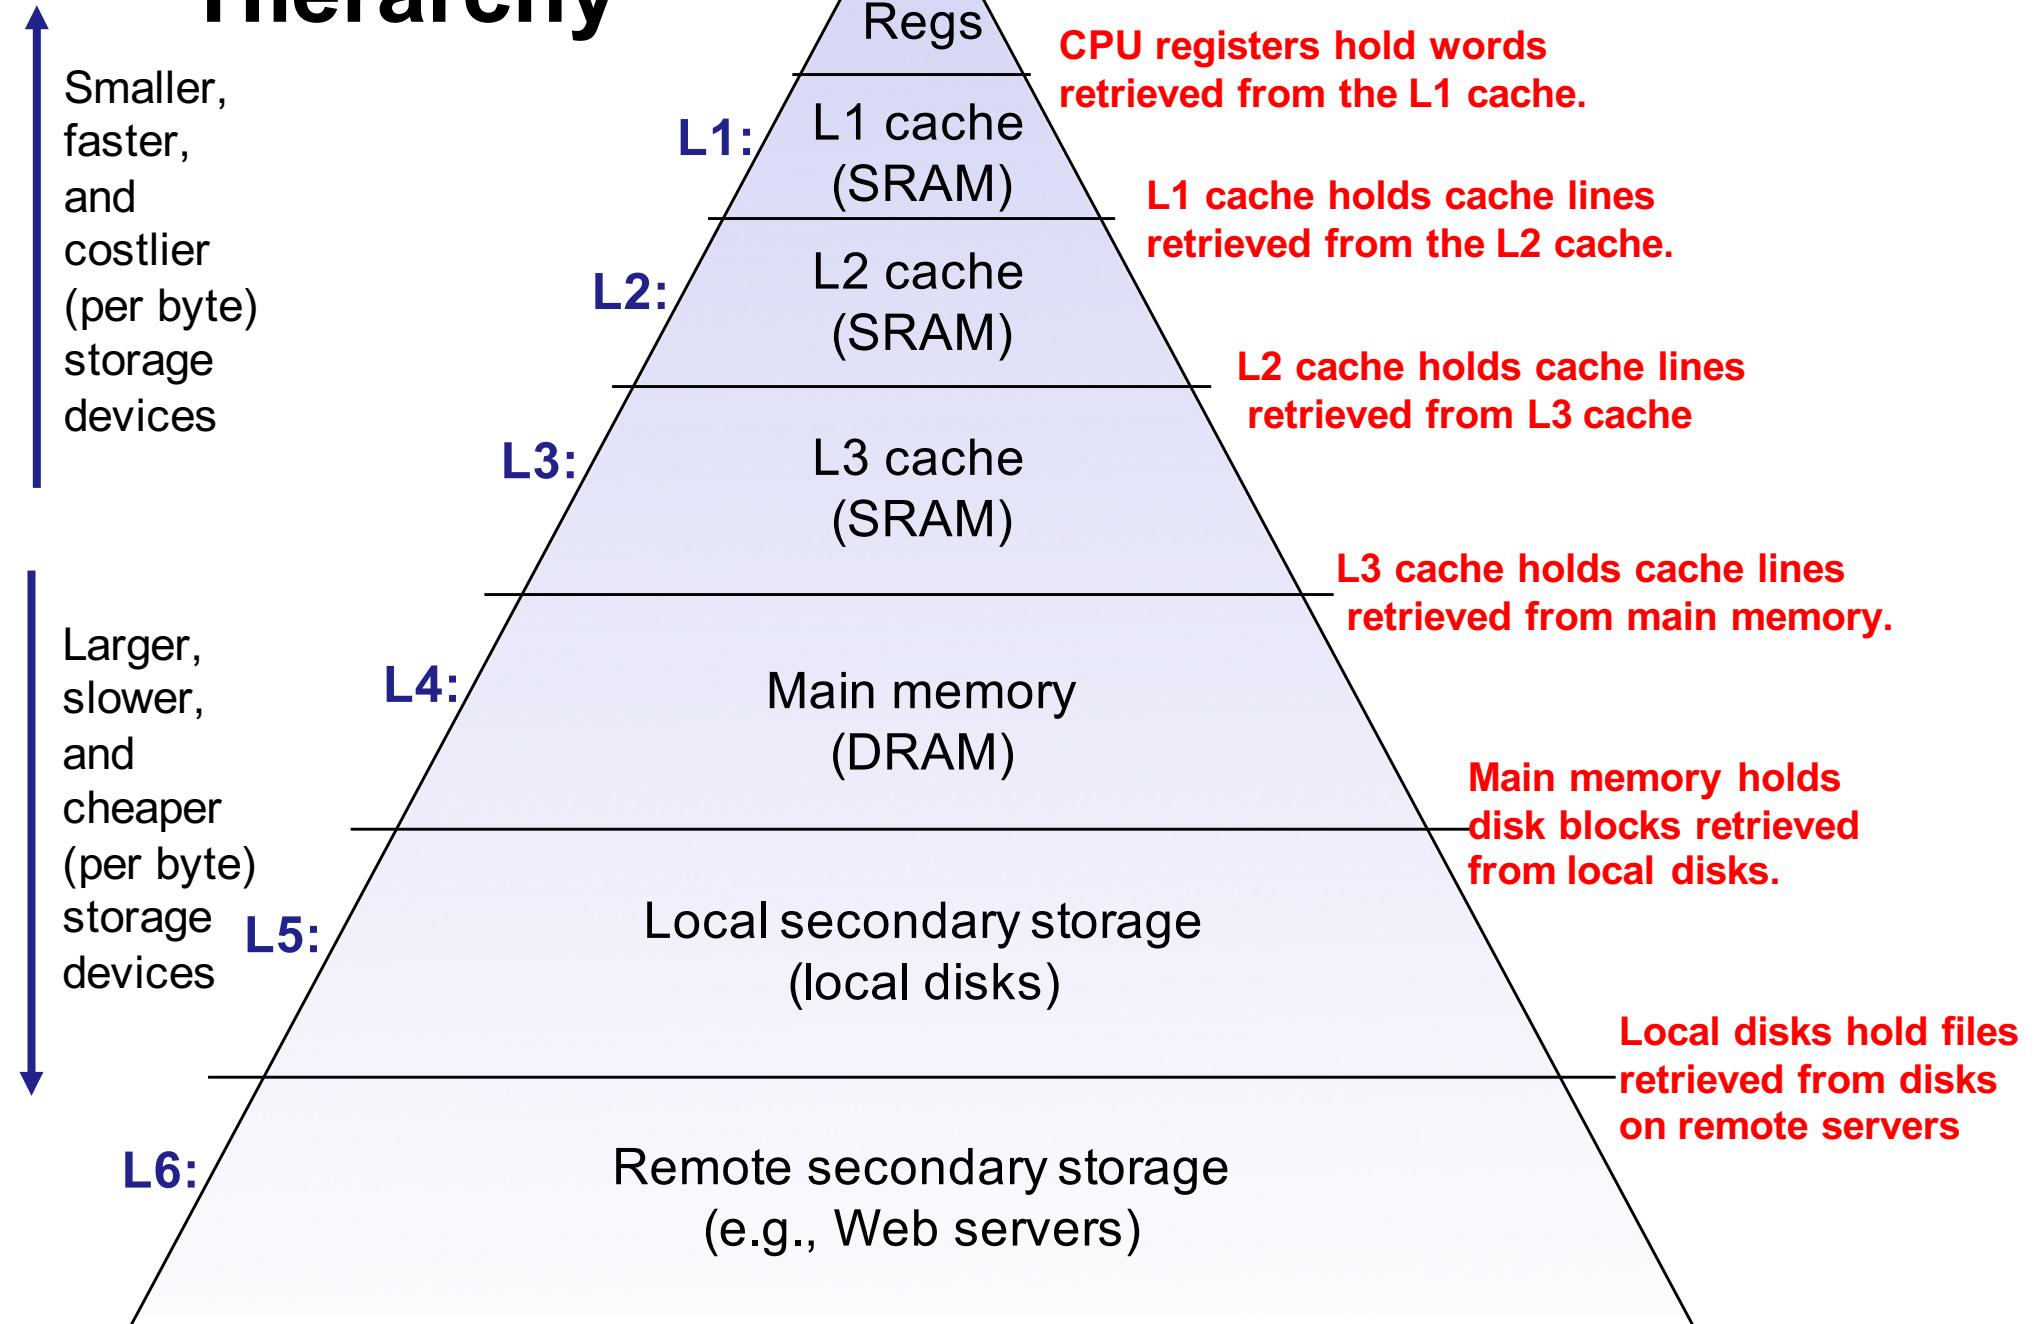
some code

```
for (i=0; i<SIZE; i++)
    sum+=A[i]
```

i, sum good temporal locality

array elements bad temporal, good spatial

Example Memory Hierarchy



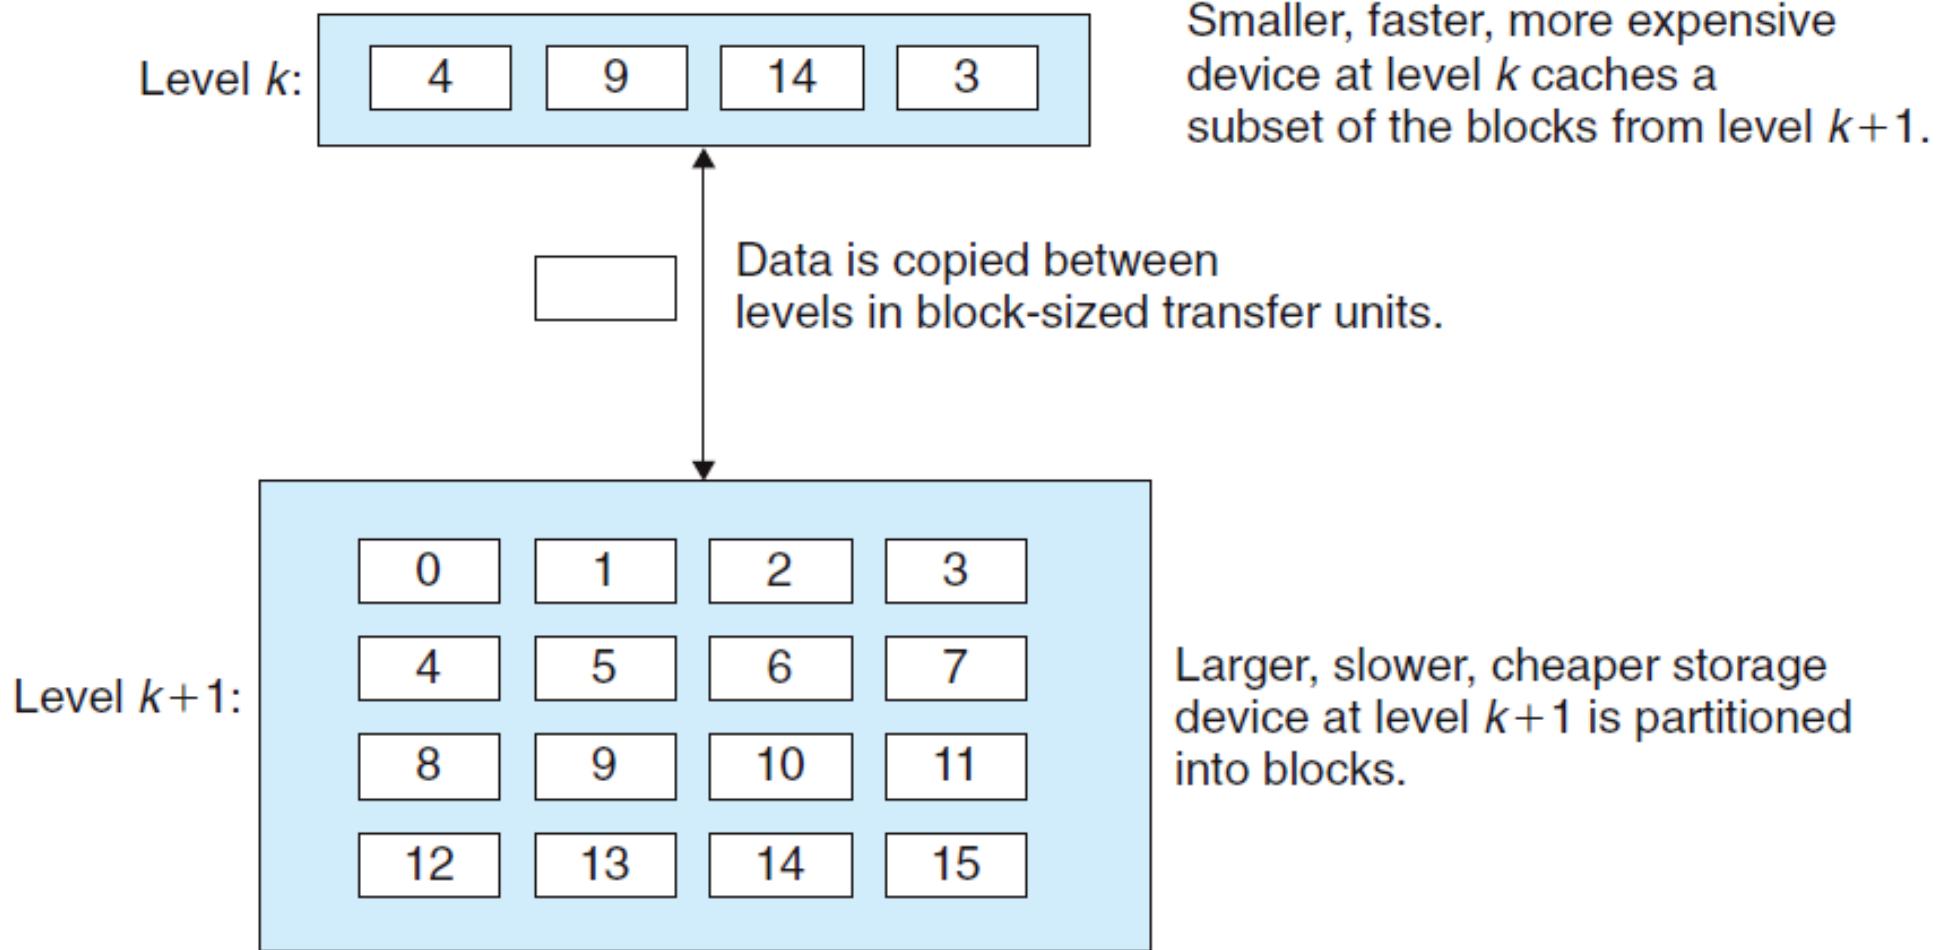


Figure 6.24 The basic principle of caching in a memory hierarchy.

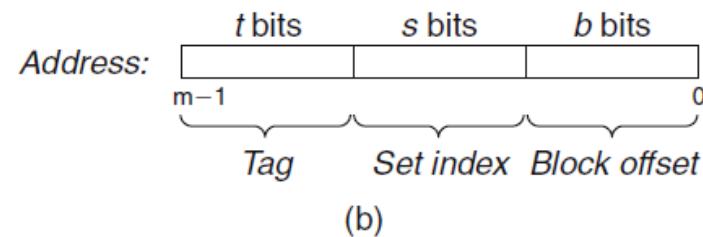
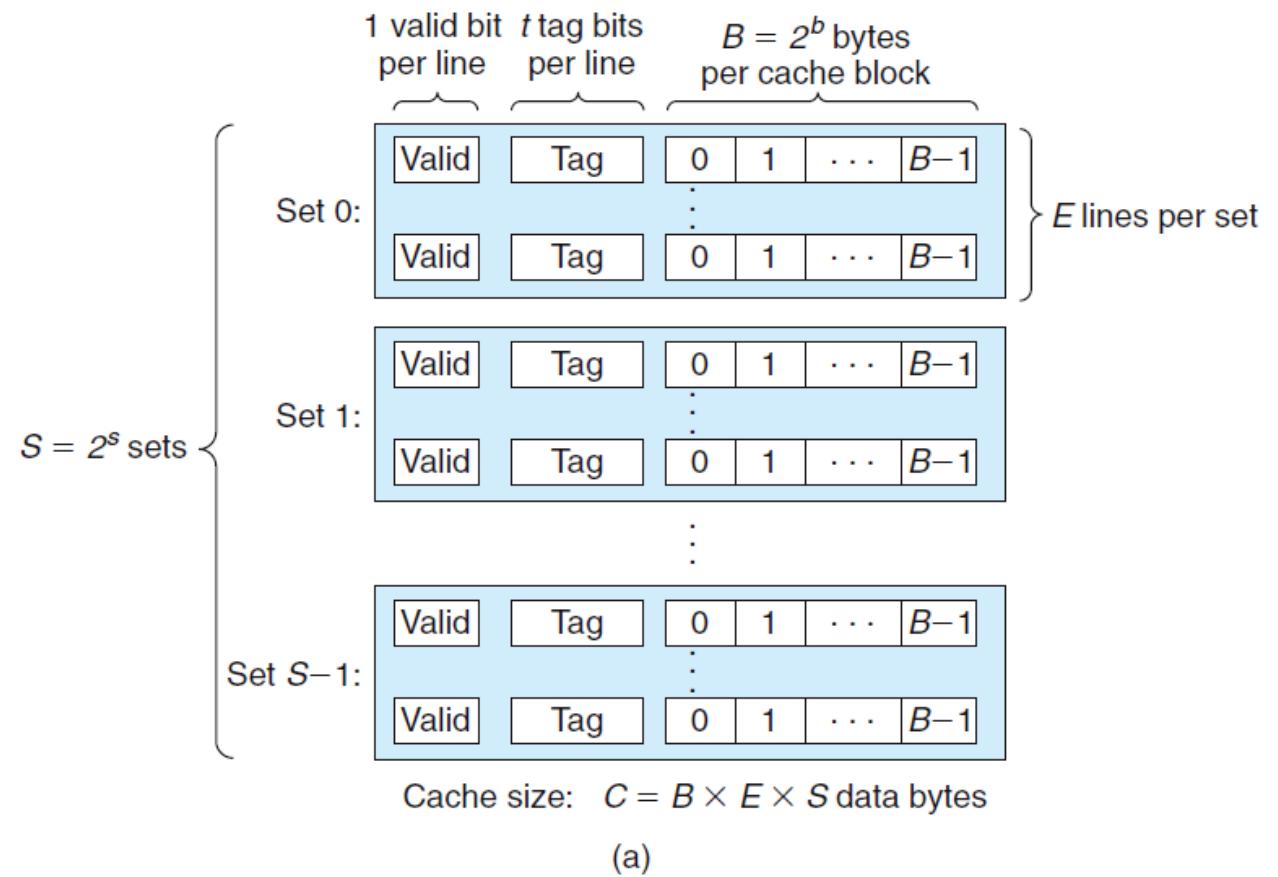
Type	What cached	Where cached	Latency (cycles)	Managed by
CPU registers	4-byte or 8-byte word	On-chip CPU registers	0	Compiler
TLB	Address translations	On-chip TLB	0	Hardware MMU
L1 cache	64-byte block	On-chip L1 cache	1	Hardware
L2 cache	64-byte block	On/off-chip L2 cache	10	Hardware
L3 cache	64-byte block	On/off-chip L3 cache	30	Hardware
Virtual memory	4-KB page	Main memory	100	Hardware + OS
Buffer cache	Parts of files	Main memory	100	OS
Disk cache	Disk sectors	Disk controller	100,000	Controller firmware
Network cache	Parts of files	Local disk	10,000,000	AFS/NFS client
Browser cache	Web pages	Local disk	10,000,000	Web browser
Web cache	Web pages	Remote server disks	1,000,000,000	Web proxy server

Figure 6.25 The ubiquity of caching in modern computer systems. Acronyms: TLB: translation lookaside buffer, MMU: memory management unit, OS: operating system, AFS: Andrew File System, NFS: Network File System.

Figure 6.27

General organization of cache (S, E, B, m).

(a) A cache is an array of sets. Each set contains one or more lines. Each line contains a valid bit, some tag bits, and a block of data. (b) The cache organization induces a partition of the m address bits into t tag bits, s set index bits, and b block offset bits.



Fundamental parameters	
Parameter	Description
$S = 2^s$	Number of sets
E	Number of lines per set
$B = 2^b$	Block size (bytes)
$m = \log_2(M)$	Number of physical (main memory) address bits

Derived quantities	
Parameter	Description
$M = 2^m$	Maximum number of unique memory addresses
$s = \log_2(S)$	Number of <i>set index bits</i>
$b = \log_2(B)$	Number of <i>block offset bits</i>
$t = m - (s + b)$	Number of <i>tag bits</i>
$C = B \times E \times S$	Cache size (bytes) not including overhead such as the valid and tag bits

Figure 6.28 Summary of cache parameters.

Practice Problem 6.10

The following table gives the parameters for a number of different caches. For each cache, determine the number of cache sets (S), tag bits (t), set index bits (s), and block offset bits (b).

Cache	m	C	B	E	S	t	s	b
1.	32	1024	4	1	_____	_____	_____	_____
2.	32	1024	8	4	_____	_____	_____	_____
3.	32	1024	32	32	_____	_____	_____	_____

Practice Problem 6.10

The following table gives the parameters for a number of different caches. For each cache, determine the number of cache sets (S), tag bits (t), set index bits (s), and block offset bits (b).

Cache	m	C	B	E	S	t	s	b
1.	32	1024	4	1	_____	_____	_____	_____
2.	32	1024	8	4	_____	_____	_____	_____
3.	32	1024	32	32	_____	_____	_____	_____

The solution is a straightforward application of the definitions of the various cache parameters in Figure 6.28. Not very exciting, but you need to understand how the cache organization induces these partitions in the address bits before you can really understand how caches work.

Cache	m	C	B	E	S	t	s	b
1.	32	1024	4	1	256	22	8	2
2.	32	1024	8	4	32	24	5	3
3.	32	1024	32	32	1	27	0	5